

Databases

Seminar 5

Indexes in SQL Server (I)

- index
 - structure stored on the disk, associated with a table or view
 - optimizes retrieval operations on the table / view
- great indexing
 - => fast applications
- poor indexing
 - => can slow down the DBMS

- syntax

```
CREATE [ UNIQUE ] [ CLUSTERED | NONCLUSTERED ]  
    INDEX index_name  
  
ON <object> ( column [ ASC | DESC ] [ ,...n ] )  
  
[ INCLUDE ( column_name [ ,...n ] ) ]  
  
[ WHERE <filter_predicate> ]  
  
[ WITH ( <index_option> [ ,...n ] ) ]
```

- index characteristics
 - clustered / non-clustered
 - unique / non-unique
 - search key - single-column / multicolumn
 - key columns / non-key columns
 - columns in the index - ascending / descending order
 - non-clustered indexes - full-table / filtered

- clustered / non-clustered index

- clustered index

- the data rows in the table are kept sorted, based on the values of the search key

```
CREATE CLUSTERED INDEX Index_Name  
    ON Schema_Name.Table_Name(Column)
```

- non-clustered index

- contains key values and pointers to the data in the table (heap / clustered index)

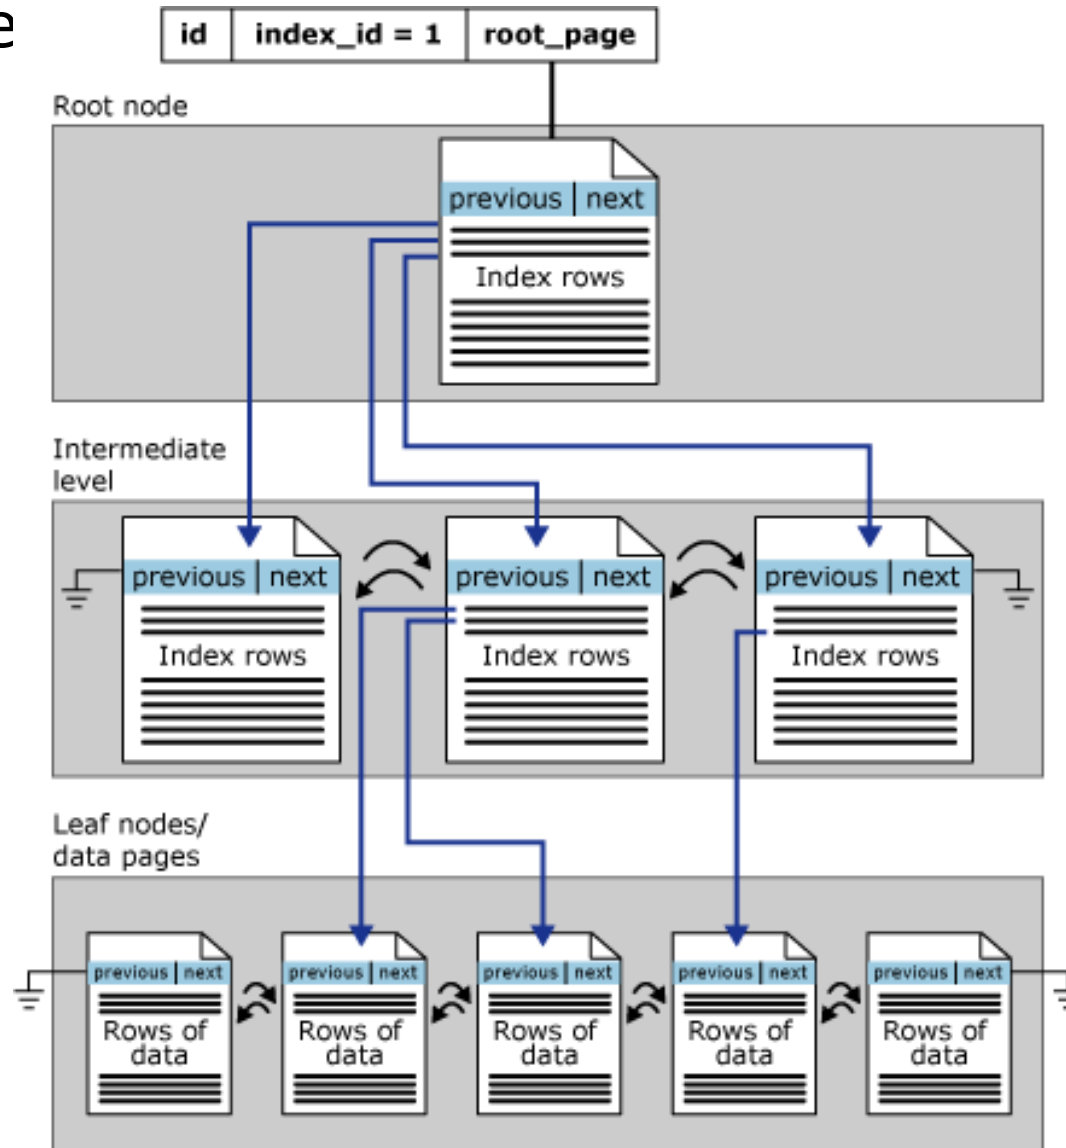
```
CREATE INDEX Index_Name  
    ON Schema_Name.Table_Name(Column)
```

- clustered / non-clustered index
 - data pages in a clustered index always include *all the columns* in the table
 - SQL Server:
 - at most one clustered index per table
 - at most 999 non-clustered indexes per table
 - an index key (clustered / non-clustered)
 - at most 16 columns, 900 bytes*

* version differences

- clustered / non-clustered index
 - clustered index
 - can be used for frequently executed queries
 - high degree of uniqueness
 - can be used in range queries
 - columns that are part of the search key:
 - shouldn't be frequently changed
 - should be narrow

- clustered / non-clustered index
 - clustered index
 - organized as a B+ tree



- clustered / non-clustered index
 - when creating a primary key on a table:
 - if a clustered index is not defined
 - a non-clustered index is not specified
- => a unique clustered index is created on the fields of the primary key

- unique indexes
 - such an index guarantees that the search key contains no duplicate values
 - specifying a unique index makes sense only when there are no entries with identical values in the key columns
 - uniqueness – useful information for the query optimizer

- key / non-key index columns
 - key columns
 - columns in the search key
 - non-key columns
 - columns specified in the INCLUDE clause when creating a non-clustered index

```
CREATE INDEX Index_Name  
    ON Schema_Name.Table_Name(Column)  
    INCLUDE (ColumnA, ColumnB, ColumnC)
```

- covering index
 - contains all the columns that are necessary in a query

- key / non-key index columns
 - non-key columns - benefits
 - columns can be accessed from the index
 - data types that are not allowed in key columns can be used in non-key columns (varchar(max), nvarchar(max), varbinary(max))
 - non-key columns are not taken into account when computing the size of the key

- filtered indexes
 - optimized non-clustered indexes
 - can be used by queries that select from a certain subset of data

```
CREATE NONCLUSTERED INDEX IDX_eid_pid_f_od  
    ON Orders (EmpId, ProdId)  
    WHERE OrderDate IS NOT NULL
```

- benefits
 - better query performance
 - reduced index:
 - maintenance cost
 - storage cost

- index design
 - analyze the characteristics of the:
 - database
 - Online Transaction Processing (OLTP)
 - Online Analytical Processing (OLAP)
 - most frequently executed queries
 - columns used in queries
 - determine the best storage location for the index

- index design - guidelines
 - database-related aspects
 - the presence of many indexes on a table deteriorates the performance of INSERT, UPDATE, DELETE, MERGE statements
 - indexing small tables is often useless

- index design - guidelines
 - query-related aspects
 - non-clustered indexes should be created on columns that are often used in WHERE and JOIN
 - covering indexes can significantly improve the performance of queries
 - as many records as possible should be changed in a single statement

- index design - guidelines
 - column-related aspects
 - length of the index key - as short as possible for clustered indexes
 - clustered indexes - better on unique / non-null columns
 - types ntext, text, image, varchar(max), nvarchar(max), varbinary(max) cannot be used for search key fields
 - column uniqueness
 - data distribution in the column
 - avoid indexes on columns with a small number of distinct values
 - filtered indexes - on columns with well-defined subsets
 - order of columns in multicolumn indexes
 - first positions - columns in equality (=), inequality (>, <, BETWEEN) conditions
 - the rest of the columns should be ordered by distinctness

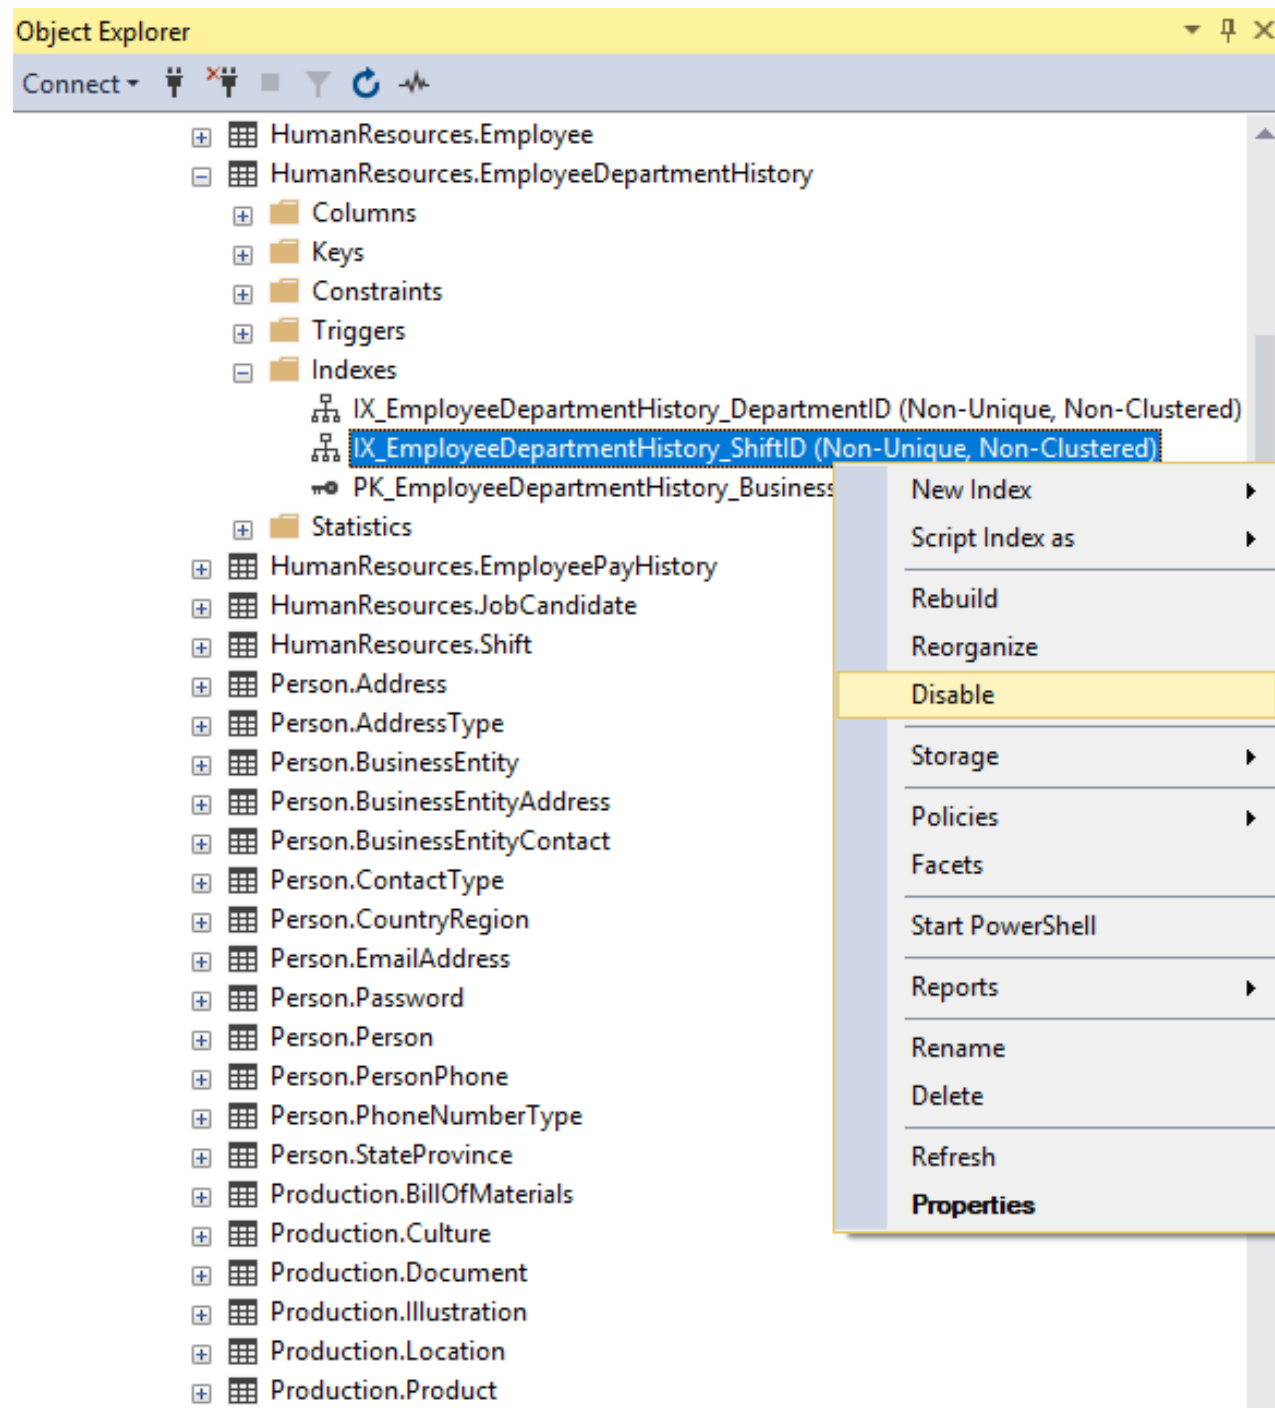
- index design - guidelines
 - column-related aspects
 - consider indexing computed columns

- indexes for deletes
 - when executing a delete statement
 - SQL Server searches for dependent rows by examining all the foreign keys (when a record r is deleted, the system checks whether r is referenced by other records)
 - if there is an index, SQL Server uses it to check the existence of related data
 - if there isn't an index, the system has to scan the identified table
 - the performance of delete operations can be improved by creating indexes on foreign keys

- **disable indexes**

```
ALTER INDEX IX_EmployeeDepartmentHistory_ShiftID  
ON HumanResources.EmployeeDepartmentHistory DISABLE
```

- disable indexes



- enable indexes

```
ALTER INDEX IX_EmployeeDepartmentHistory_ShiftID  
ON HumanResources.EmployeeDepartmentHistory REBUILD
```